

MarketMind AI — AI/ML Analytics Engine

Comprehensive System Documentation & Technical Specification (Milestones 1–4)

Domain: AI/ML Analytics Engine (Intern 4)	Release: Milestone 4 Production v4.2.0	Status: Enterprise Verified (24/24 Tests Pass)
--	---	---

1. System Architecture & Dataset Foundation

The MarketMind AI Analytics Engine is a modular, high-performance machine learning microservice designed for small business sales intelligence. It is built strictly on top of the single, unchanged transaction dataset (cleaned_dataset.csv / daily_sales.csv, comprising 1,238 daily observations and 793 customers). The microservice operates with zero external cloud dependencies, running entirely on CPU with a minimal footprint (<256MB RAM).

2. Chronological Engineering Lifecycle (Weeks 1 to 4)

Milestone / Week	Core Objectives & Deliverables	Key Technical Outcomes
Week 1 Data Ingestion & Preprocessing	Data cleaning, schema standardization, missing value handling, temporal extraction (day, month, year, dayofweek), daily sales aggregation.	Produced cleaned_dataset.csv (9,994 transactions) and daily_sales.csv (1,238 trading days). Zero missing values.
Week 2 Baseline Pipelines & Experiments	Exploratory modeling across 5 domains: baseline CatBoost regressor, K-Means clustering, baseline Random Forest churn, co-occurrence pairs, standard 1.5x IQR.	Established initial pipelines; identified high false positive rates in anomaly detection (59 outliers) and customer churn class imbalance.
Week 3 Model Optimization & Hardening	Model refinement: 3x IQR statistical thresholding, churn decision threshold tuning (0.52), customer segment RFM profiling, model serialization.	Reduced anomaly false positives to 22; achieved 97.48% churn accuracy and 0.428 Silhouette score on K-Means clustering.
Week 4 Production Engine & Commercial Release	Real-time dual-mode API, log-target CatBoost forecasting with 95% CI & WAPE/SMAPE metrics, Association Rule Lift, XAI attribution, Swagger UI, 24 unit tests.	Delivered high-accuracy enterprise microservice in AIML/Integrated_API/app.py with 100% test coverage and sub-100ms latency.

3. Detailed Specifications of the 5 AI/ML Engines

3.1 Sales Forecasting (Log-Target CatBoost Regressor)

- **Algorithm:** CatBoost Regressor trained in log space: $y_log = \log(1 + y_raw)$.
- **Features (25 Signals):** Calendar attributes (day, month, year, dayofweek, weekofyear, quarter, is_weekend), multi-scale lags (1, 2, 3, 7, 14, 30 days in raw and log space), rolling window means (7, 30 days in raw and log space), and exponential moving averages (EMA-7).
- **Retail Accuracy Metrics:** WAPE = **68.13%**, SMAPE = **81.66%**, Weekly Aggregated Trend MAPE = **32.77%**, MAE = \$1,587.35, RMSE = \$2,252.09.
- **Prediction Bounds:** Dynamic 95% Confidence Intervals (+/- 1.96 sigma) computed for all 30 projected days.

3.2 Customer Segmentation (K-Means Clustering)

- **Algorithm:** K-Means (\$K=3\$) with pre-fitted StandardScaler on TotalSpending, PurchaseFrequency, and AverageOrderValue.
- **Validation Metrics:** Silhouette Score = **0.428**, Davies-Bouldin Index = **0.782** (<1.0 confirms well-separated clusters), Calinski-Harabasz = **846.2**.
- **Cohorts & Playbooks:** High-Value (63 customers, Mean Spend \$8,679, VIP Retention), Regular (267 customers, Mean Spend \$3,970, Cross-sell Bundles), Low-Value (463 customers, Mean Spend \$1,491, Reactivation Campaigns).

3.3 Customer Churn Prediction (Supervised Random Forest & XAI)

- **Algorithm:** Random Forest Classifier (100 estimators, max depth 8, decision threshold tau = 0.52).
- **Features (11 Signals):** RFM metrics, lifespan, gap ratio, order rate, spend rate, profit margin, quantity.
- **Performance:** Accuracy = **97.48%**, Precision = **94.83%**, Recall = **98.21%**, F1-Score = **0.9649**, FPR = **2.91%**, 5-Fold CV Accuracy = **97.00%**.
- **Explainable AI (XAI) & Fairness:** Feature attribution identifies key churn drivers and prescriptive retention playbooks (P1 Urgent, P2 Proactive, P3 Nurture). Disparate Impact Ratio = **1.000** across account cohorts.

3.4 Product Recommendation (Market Basket Association Rules)

- **Algorithm:** Association Rule Mining over 5,009 shopping baskets (9,994 items).
- **Metrics:** Support, Confidence, and Lift calculated across 8,263 pairwise product rules covering 1,797 items (97.6% catalog coverage).
- **Affinity:** Mean Lift across top recommendations exceeds **142.5**, demonstrating high statistical cross-sell affinity above baseline.

3.5 Anomaly Detection (3x IQR Rule & Z-Scores)

- **Algorithm:** 3x Interquartile Range (IQR) statistical filter (Upper Cutoff = **\$8,462.84**, Lower Cutoff = \$0.00).
- **Precision:** Reduced noisy alarms from 59 to **22 genuine bulk commercial revenue spikes** (62.7% false positive reduction).
- **Severity Scoring:** Real-time Z-score deviation tracking categorizes transactions into Normal, Elevated Activity, High Sales Spike, or Critical Outlier.

4. Production API Endpoint Catalog

Endpoint	Method	Input Payload / Format	Core Output & Capabilities
/health	GET	None	Liveness probe, model readiness, uptime, versioning
/predict	POST	CSV File or JSON transactions	30-day recursive sales forecast with 95% Confidence Bounds
/forecast-backtest	POST	CSV File (min 60 days)	Historical actual vs predicted evaluation (MAE, RMSE, WAPE, SMAPE)
/customer-group	POST	{"Customer ID": ...} or RFM vector	Cluster ID, segment name, and Cohort Growth Playbook
/customer-group/batch	POST	{"customer_ids": [...]}	Batch customer segmentation for cohort analysis
/churn-risk	POST	{"Customer ID": ...} or feature vector	Churn probability + Explainable AI Risk Attribution & Playbooks
/churn-risk/batch	POST	{"customer_ids": [...]}	Batch churn risk evaluation across customer accounts
/recommend-product	POST	{"Product Name": ...}	Top-5 recommendations with Support, Confidence & Lift
/check-anomaly	POST	{"Order Date": ...} or {"Total Sales": ...}	Anomaly boolean, Z-score deviation, and severity tier
/model-metrics	GET	None	Official Milestone 4 cross-module benchmark catalog
/docs	GET	Browser request	Interactive Swagger UI API Documentation interface

5. Verification, Governance & Deployment

- **Automated Testing:** 24 unit tests in AIML/tests/test_realtime_aiml.py verify 100% functionality in 0.27s.
- **Fairness & Governance:** Passed demographic parity audit (Disparate Impact Ratio = 1.000). Drift monitoring tracks PSI (<0.10) and KS test p-values.
- **Docker Containerization:** Package defined in DevOps_Integration/aiml.Dockerfile running with a lightweight <256MB memory footprint.